

```
> init_evaluation_suite()
```

The Classification Diagnostic Playbook

Translating mathematical evaluation
metrics into practical business
decisions.

Classification models must balance predictive power with inferential trustworthiness.

The Goal: Prediction



Classification predicts categorical values, assigning incoming data into predefined buckets (e.g., churn vs. no churn).


The Reality: Inference



A model must not just be accurate, but understandable. A highly inaccurate model won't be trusted for inference, and understanding variable impact improves real-world accuracy.

You often need to strike a balance between these goals – a model that is very inaccurate won't be too trustworthy for inference.


Rigorous evaluation requires isolating data into distinct testing environments to prevent overfitting.



Different classification algorithms generate predictions using distinct mathematical mechanisms.

The Classification Engine Room			
Algorithm	Core Mechanism	Key Hyperparameters	Watch Out For
Decision Tree	Splits data based on feature values at each node.	max_depth	Prone to overfitting if not pruned properly.
Linear SVM	Finds optimal hyperplane separating classes with maximum margin.	C, kernel	Sensitive to noise and extremely large datasets.
K-Nearest Neighbors	Classifies based on the majority class of nearest historical data points.	n_neighbors	Computationally expensive as the dataset grows.
XGBoost	Gradient boosting builds trees sequentially to correct previous errors.	n_estimators, learning_rate	Complex to tune and computationally intensive.

Distance-based classifiers define spatial boundaries based on the proximity of historical data points.




Spatial Mechanics: K-Nearest Neighbors

K-Nearest Neighbors (KNN) classifies unknown observations based on the characteristics of nearby labeled points.

Optimization Strategy

Test various 'k' values (e.g., `n_neighbors=5`) and measure accuracy while considering class distribution and feature relevance.

Ensemble methods reduce individual model variance by aggregating multiple weak predictions into one robust final prediction.



Random Forests & The Bias-Variance Tradeoff

Random forests use “bagging” to train multiple decision trees on bootstrapped data subsets. By aggregating the results, the model drastically improves accuracy and reduces the variance typical of individual decision trees.

```
model = RandomForestRegressor(n_estimators=100, max_depth=5)
```


The Confusion Matrix is the foundational diagnostic tool that summarizes predicted versus actual classes.

Confusion Matrix Decoder

PREDICTED CLASS (0, 1)	
ACTUAL CLASS (0, 1)	0
	1
0	True Negative (TN)
1	False Positive (FP) - Type I Error
0	False Negative (FN) - Type II Error
1	True Positive (TP)

This is the most common and concise way to evaluate performance. All other core metrics—accuracy, precision, and recall—are derived directly from these four quadrants.

Evaluation metrics are targeted ratios extracted from specific quadrants of the confusion matrix.

Accuracy

% of total classified correctly.

$$\frac{(TP + TN)}{\text{Total}} = \frac{\begin{array}{|c|c|} \hline TP & TN \\ \hline \end{array}}{\begin{array}{|c|c|} \hline TN & FP \\ \hline FP & TP \\ \hline \end{array}}$$

Precision

% of predicted positives classified correctly.

$$\frac{TP}{(TP + FP)} = \frac{\begin{array}{|c|} \hline TP \\ \hline \end{array}}{\begin{array}{|c|c|} \hline FP & TP \\ \hline \end{array}}$$

Recall

% of actual positives classified correctly.


$$\frac{TP}{(TP + FN)} = \frac{\begin{array}{|c|} \hline FN \\ \hline \end{array}}{\begin{array}{|c|} \hline TP \\ \hline \end{array}}$$

Selecting the right metric depends entirely on the business cost of false positives versus false negatives.

The Ultimate Metric Decision Matrix

Metric	Business Translation	Formula	When to Optimize
Accuracy	Is the overall model correct?	$(TP+TN) / Total$	Use when data is balanced and both outcomes carry equal risk.
Precision	When it says YES, can I trust it?	$TP / (TP+FP)$	Use when False Positives are a major risk (e.g., spam filters).
Recall	Did it miss any real targets?	$TP / (TP+FN)$	Use when False Negatives are a major risk (e.g., fraud detection).
F1 Score	Is it balanced?	Harmonic Mean of Precision & Recall	Use for imbalanced data requiring a balance between FP and FN.
ROC AUC	Does it rank probabilities well?	Area Under Curve	Use to evaluate how well predicted probabilities rank positive instances above negative ones.

Soft classification models predict probabilities, requiring a decision threshold to assign the final binary label.



Threshold Shifting

By default, models use 0.5 to decide between True and False.


Adjusting this slider changes the balance of False Positives and False Negatives, directly impacting Precision and Recall.

Terminal

```
fpr, tpr, thresholds =  
sklearn.metrics.roc_curve(y_true,  
y_prob)  
auc_score = sklearn.metrics.auc(  
fpr, tpr)
```


Multi-class classification extends binary logic using either One-vs-Rest or One-vs-One strategies.

One vs Rest (OVR)



Mechanism: Trains one classifier per class, distinguishing a single class against all others combined.

Pros


Simpler, highly scalable.

Cons

Less accurate for highly imbalanced classes.

```
model = OneVsRestClassifier(LogisticRegression())
```

One vs One (OVO)



Mechanism: Trains a distinct classifier for every possible pair of classes in the dataset.

Pros

Highly effective for small datasets.

Cons

Computationally expensive as the number of classes grows.

```
model = OneVsOneClassifier(LogisticRegression())
```


Scoring multi-class matrices requires isolating each class to calculate its unique true and false outcomes.

	Predicted Product A	Predicted Product B	Predicted Product C	Predicted Product D
Actual Product A	214	1	8	2
Actual Product B	15	0	0	0
Actual Product C	0	0	0	0
Actual Product D	3	0	0	0

Isolating Product A

True Positives:
214 (Predicted A, Actual A)


Precision:
 $214 / (214 + 15 + 3) = .9224$

Recall:
 $214 / (214 + 1 + 8 + 2) = .9511$

The True Negatives include all cases where Product A was neither predicted nor observed (the grayed-out cells).

Multi-class performance is summarized by taking an average of the metrics across all individual classes.

	PREDICTED			
	A	B	C	D
	214	1	8	2
	15	0	0	0
	0	0	0	0
	3	0	0	0
ACTUAL	A	B	C	D




`average='macro'`: A simple, unweighted average of the metric across all classes.


`average='weighted'`: Returns an average based on class frequency. Ideal for imbalanced datasets!

```
precision_score(y_actual, y_pred, average='weighted')
```


If you use the wrong metric to guide the modeling process, it doesn't matter which model you use.



Understanding how and when to use classification metrics is arguably as important as knowing the algorithms themselves.
Business reality must dictate the math.



1. Define the business cost of failure (False Positives vs. False Negatives).



2. Optimize the model for 1-2 core metrics (e.g., Precision vs. Recall).



3. Use all remaining metrics purely for context when evaluating final validation performance.