


Beyond the Agent Hype: Architecting Predictive ML and GenAI in Production


A strategic playbook for designing durable, closed-loop AI systems



The most useful mental model is not “two kinds of AI glued together.”

Prediction accuracy alone does not guarantee good business decisions. True autonomy requires separating estimation from recommendation.

The “Two Brains” Engine




Distinct and complementary roles inside the decision system.

Neither Machine Learning nor Generative AI can safely solve enterprise problems alone.

	Predictive ML	Generative AI	Optimization / Rules
Superpower	Learning repeatable relationships from historical data.	Synthesizing heterogeneous context and conversational interaction.	Enforcing compliance, budgets, capacities, and deterministic invariants.
Fatal Flaw	Correlation is not causation; produces a score without operational context.	Confabulation; omits constraints; persuasive but unsupported reasoning.	Brittle for ambiguous or novel situations.
Output	Probability, forecast interval.	Recommendation narrative, rationale.	Optimal action set; Allow/Deny.


The End-to-End Blueprint: From Signal to Autonomous Action

Prescribe actions based on predictions, validate them deterministically, and capture outcomes to close the learning loop.




The Immutable Boundary: The Prediction Contract.

The GenAI agent receives a read-only serialization. It may explain it, but it must not recalculate or overwrite prediction fields.




Building the Context: Injecting reality before generation

Generative AI requires a complete state package—predictions, knowledge, and operational limits—to draft a feasible response.



Deterministic Guardrails: The Red Filter

A GenAI-generated recommendation is not automatically a valid prescription. Every action must pass through deterministic validation.




"OWASP identifies excessive functionality and autonomy as root causes of agent risk. Enforce least privilege and downstream authorization."

The Dual-Feedback Loop: Severing self-reinforcing cycles.

Evaluate predictive accuracy separately from prescription utility.




Do not automatically retrain the ML model using GenAI recommendations.






The Architecture in Practice: Cross-Domain Modularity

The exact same architectural boundary applies whether predicting IT outages or financial variances.

IT Operations



Finance / Scenario




Agent and Workflow Framework Selection

The engine you choose depends on your center of gravity: durable state, RAG, or Microsoft integration

	Strengths	Verdict
Temporal / Camunda	Durable state, timers, approvals, end-to-end audit	Production default for Phases 4–5 (Human-Approved/Autonomous)
LangGraph	Low-level state graphs, deterministic/agentic mixing, persistence	Strong cloud-neutral default for Phase 3 (Recommendation Agent)
Microsoft Agent Framework	Typed tools, telemetry, MCP, deep .NET integration	Strong enterprise option with disciplined versioning
LlamaIndex Workflows	Event-driven, async, tightly coupled with ingestion	Best when RAG/document retrieval is the center of gravity


Orchestration Constraints: Bounding the Agent.

Keep CrewAI in the role it is best suited for: orchestration, interpretation, and synthesis.
Do not let agents execute arbitrary math.




The 5-Phase Maturity Roadmap.

Preserve stable, typed service boundaries while progressively adding orchestration, policy, and recovery controls.



The Autonomous Illusion

Moving to an 'Autonomous Agent' does not mean deploying a massive, unconstrained LLM. It is an incremental policy change over a well-controlled Phase 4 architecture.



Every action remains policy-checked, identity-bound, idempotent, observable, and rollback-capable. NIST and OWASP emphasize least privilege.

The Golden Rules of the Prediction-to-Prescription Pipeline.

Rule of thumb: Use deterministic code for known transformations, a workflow engine for known processes, an optimization engine for constrained choices, and an LLM agent only for interpreting ambiguity.

1. Typed Boundaries

The LLM receives a frozen, immutable prediction envelope. It explains, it does not calculate.

2. Deterministic Guardrails

Validate schemas, permissions, and policy compliance in standard application code before any execution.

3. Durable Orchestration

Never delegate high-impact durability or authorization solely to a probabilistic agent loop.

4. Split Evaluation

Measure predictive ML (accuracy, drift) strictly separate from Generative AI (utility, groundedness).